



# COSI 127B: Introduction to Database Systems

## Programming Assignment II

*Database Management and Applications*

**Due: 3:30 p.m., Monday, March 20, 2017**

### 1 Introduction

The purpose of PA1 was to let you get familiar with SQL by composing several complex queries. PA2 is a complementary assignment designed to let you get familiar with three important skills required of a DBA: writing integrity constraints, implementing triggers, and integrating SQL with Java. For this assignment, you will act as DBA for a manufacturing company that must monitor its inventory of parts, the suppliers for those parts, orders made for those parts and their fulfillments. You will be provided with the schema and data to populate the very large database maintained by the manufacturing company. Your job will be twofold:

1. To formulate and construct integrity constraints and triggers to ensure data integrity within this large application.
2. To explore the use of programmatic database driver to user facing power applications.

### 2 Background

The data you will be working with is drawn from the TPC-H benchmark<sup>1</sup> database, which models the data needs of a manufacturing company. Information maintained in this database includes an inventory of parts, suppliers for those parts, orders for parts and order fulfillments. Your job will be to tune this database by adding integrity constraints and triggers, as well as a Java-based user-interface to facilitate certain interactions.

#### 2.1 The TPC-H Database

The database you will be working with consists of eight tables, each with several attributes. Each attribute name is prefixed by a unique prefix (derived from the name of the table where it is contained), to ensure the uniqueness of attribute names across the database. Throughout this document, we will refer to an attribute without its prefixes if

---

<sup>1</sup>A benchmark consists of a schema, data population and a suite of queries and transactions that is used to compare the performance of different database systems.

the table containing the attribute is understood. The eight tables, their primary keys, and their assumed attribute prefixes are shown in on the following page.

### 2.1.1 Tables

Table Name	Prefix	Primary Key
region	r	regionkey
nation	n	nationkey
supplier	s	suppkey
part	p	partkey
partsupp	ps	partkey, suppkey
customer	c	custkey
lineitem	l	orderkey, linenumber
orders	o	orderkey

Table 1: The tables of the TPC-H Database

## 2.1.2 Schema

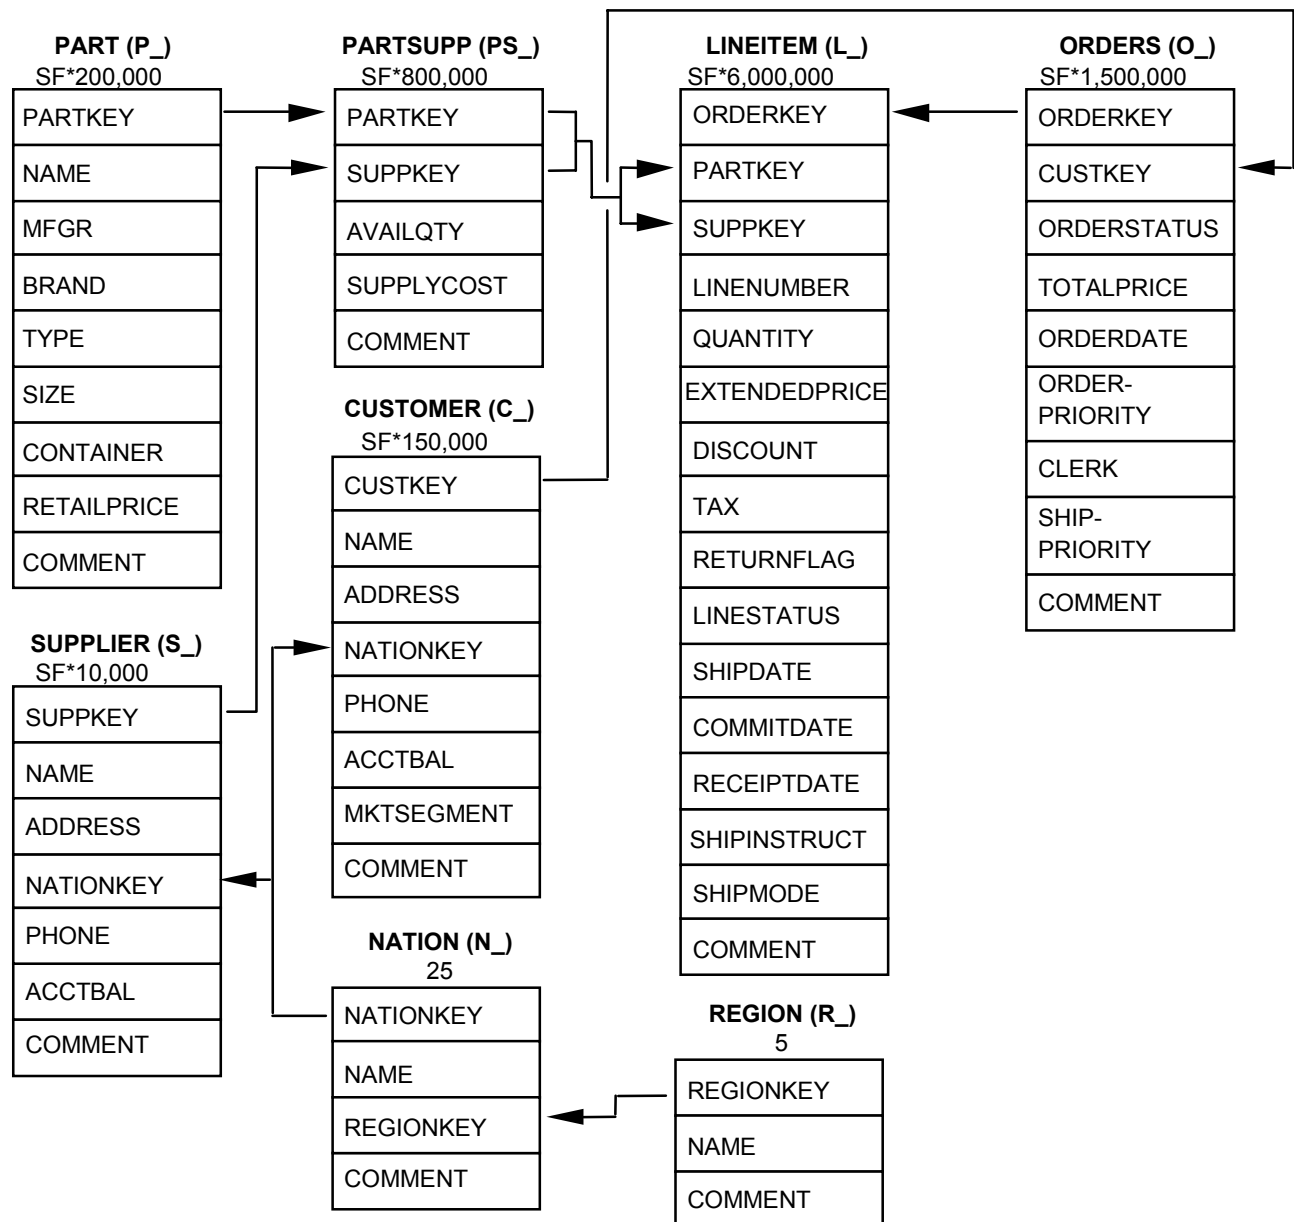


Figure 1: The Schema of the TPC-H Database

### 2.1.3 Description

**region:** This table lists the 5 regions (i.e., continents) in the world the data in this database pertains to. Each region has a name (name) and a comment permitting annotations for the region (comment).

**nation:** This table lists all nations the manufacturing company deals with. Every nation has a name (name), a region in which it is contained (regionkey) and a comment (comment).

**supplier:** This table lists all suppliers of interest for the manufacturing company. Data for a supplier includes its name (name), address (address), the nation where it is located (nationkey), phone (phone), account balance (acctbal) and comments about the supplier (comment).

**part:** This table holds information about all parts needed by the manufacturing company. Relevant data for a part includes its name (name), the parts manufacturer (mfr), and its retail price (retailprice).

**partsupp:** A row in this table represents a given part (partkey) that is supplied to the manufacturer by a given supplier (suppkey). For each part and supplier, additional data includes the number of parts available from the supplier (availqty) and the suppliers cost for the part (supplycost).

**customer:** This table contains information regarding the customers of the manufacturing company. Information stored for each customer includes his/her name (name), address (address) and nation he/she resides in (nationkey), as well as a contact phone number (phone), and the balance of the customers account with the company (acctbal).

**lineitem:** This table maintains each individual part order contained in a purchase order. Every row represents some ordered part (identified by the part (partkey) and the supplier it was ordered from (suppkey)) the quantity ordered (quantity), the total price paid for the parts (extendedprice) before tax and a discount, the date when the parts were shipped (shipdate), the date when the parts were promised by the supplier (commitdate), the date when the parts were received (receiptdate), the discount received (discount) and so on. As well, each lineitem contains an attribute (linestatus) that has a value of CO if the part order is still open and CF if the part order has been fulfilled.

**order:** This table contains details about every purchase order (or just, order). An order has a single customer (though a customer can maintain multiple orders) and is made up of one or more lineitems. An order status (orderstatus) is set to CO if it is open (i.e., if all lineitems it contains have an open status), CF if the order is fulfilled (i.e., every lineitem it contains is fulfilled), and CP if it is partially fulfilled (i.e., if some but not all lineitems in the order have been fulfilled). Order records also include the total price of the order (totalprice) and the date the order was placed (orderdate).

## 3 Getting Started

This section will instruct you on how to setup your development environment where you can write your queries for part 1 and your code for part 2.

Please proceed in order through the following sections, and refer back to them as needed throughout the assignment.

### 3.1 Environment

All queries and Java code written as part of this assignment will be executed on one of the CS public machines. . They are accessible via SSH and are able to connect to the database servers via **PSQL**. The individual server you will use are the same as the one you used for PA1.

To begin, simply log onto the server using **SSH**:

```
$ ssh [your username]@[cs public machine name].cs.brandies.edu
```

Once logged in to the server, you will need to download all the code needed for the assignment. You will copy an archive containing this code from the class directory at `/home/o/class/cs127b/PA2`. To do so, simply execute the following:

```
$ cp /home/o/class/cs127b/PA2/PA2Files.tar.gz .
```

Then proceed to decompress and untar the archive using `tar`:

```
$ tar -xvzf PA2Files.tar.gz
```

This will create a directory called “MyPA2” containing all of the files you will need for this assignment. This will be your working directory for the remainder of this guide.

## 3.2 Database

To log onto your database you will need to set up your environment variables. These will be read by `PSQL` and `JDBC` to authenticate and log you in to the right database. Simply execute the following set of commands in your shell<sup>2</sup>:

```
$ export PGHOST=[your database host]
$ export PGUSER=[your user name]
$ export PGDATABASE=[your user name]pa2
$ export PGPASSWORD=[your database password]
```

### 3.2.1 Loading TPC-H

Once you set up these variables, you will want to populate your database with TPC-H data. This task along with many other tasks in this assignment will be automated using `make` - a `*nix` tool to execute a predefined batch of commands.

To create a new database and load it with TPC-H data, execute the following command:

```
$ make initialize-db
```

This should take about a minute. Once it is done, enter `psql` and type `\d` to show all tables. You should see the following output:

```

              List of relations
Schema |   Name   | Type  | Owner
-----+-----+-----+-----
public | customer | table | [your user name]
public | lineitem | table | [your user name]
public | nation   | table | [your user name]
public | orders   | table | [your user name]
public | part      | table | [your user name]
public | partsupp | table | [your user name]
public | region   | table | [your user name]
public | supplier | table | [your user name]
(8 rows)
```

Ensure the data is there by selecting from one of the tables.

### 3.2.2 Database Reset

If ever you feel like you’ve corrupted the database by doing an exercise incorrectly, you can always start again with a fresh, uncorrupted database by running the following reset command:

```
$ make reset-db.
```

The above should also be used for several exercises in this assignment, which require that you reset to a fresh database so as to erase the effects of previously completed exercises.

---

<sup>2</sup>You can also add these to your `.bashrc` to make these variables persistent.

### 3.3 JDBC

The second part of the assignment will involve writing both SQL code to fetch a result set from the database using JDBC and Java code to process the result set into a specific output using the provided API's. Template code to do this as well as the necessary libraries are provided.

#### 3.3.1 Provided Materials

Included in the assignment directory are the following components:

- **makefile** - A makefile to ease setup and compilation
- **README** - this list of files
- **lib/jdbc.jar** - The Postgres database driver
- **lib/pa2.jar** - The collection of API's for Gnuplot, Graphviz, and LaTeX
- **docs/index.html** - Javadoc documentation for said API's
- **src/part1.java** - Template code for economic growth plot (part 2.1)
- **src/part2.java** - Template code for global trade graph (part 2.2)
- **src/part3.java** - Template code for purchase orders (part 2.3)
- **txt/orders.txt** - Text file containing purchase orders (needed for part 2.3)

There are many initially empty directories as well, such as **gv**, **plot**, **tex**, **bin**, and others. Those will be filled by the output from your code. The **sql** directory is for containing your code from Part 1.

Look at each file, and **read the documentation**.<sup>3</sup>

Note that the provided Java code are given for your convenience. Feel free to write your own code from scratch - at your own risk. All you Java program needs to do is accept no arguments and print the resulting **gnuplot**, **graphviz**, or **L<sup>A</sup>T<sub>E</sub>X** output to the console (via **System.out.println()**). Note that for part 3 your program also needs to accept orders, one per line, from standard in.

If you do plan on writing your own code from scratch, be aware all JDBC connection parameters (**PGDATABASE**, **PGHOST**, **PGUSER**, and **PGPASSWORD**) must be fetched from environment variables, as is shown in the provided code. Do not hard code your own connection parameters (it makes it hard to test your code, and is bad practice anyways). Use **System.getenv()** in Java to get an environment variable from the current session.

#### 3.3.2 Setup

Once all the files are on the server and you feel familiar with the basic API's, you can get started.

The included **makefile** will again help you compile and run your program. All it does is link your program to the necessary JAR's - inspect it for more information about how it works.

To build your program, simply run:

```
$make build-[part1 / part2 / part3]
```

The resulting classfiles will output to **./bin**.

To run your program, execute the following command:

```
$make run-[part1 / part2 / part3]
```

The output will be directed to standard out and redirected to either **plot**, **gv**, or **tex** depending on the part you are running (part1, part2, or part3 respectively).

To produce a visualization from your program, execute the following command:

---

<sup>3</sup>Documentation for JDBC can be found at <http://docs.oracle.com/javase/7/docs/api/java/sql/package-summary.html>

```
$make vis-[part1 / part2 / part3]
```

The visualization will be saved in `./vis` using input from `plot`, `gv`, or `tex` depending on the part you are visualizing (part1, part2, or part3 respectively).

## 4 Constraints and Triggers

### 4.1 Introduction

Your task for this part of the assignment is to add integrity constraints and triggers to the TPC-H database to ensure that data consistency is maintained. Although triggers have not been discussed in class, you can read about triggers at:

<http://www.postgresql.org/docs/8.4/static/sql-createtrigger.html>

or

<http://www.postgresql.org/docs/8.4/static/plpgsql-trigger.html>

In short, triggers are statements that a database executes automatically as a side effect of certain updates to the database. When registering a trigger with the system, you specify what events (expressed as conditions) trigger execution of the trigger, and what actions (expressed in the Postgres trigger language) are to take place as a result. For example, a trigger could be used to automatically calculate the value of derived attributes, every time one of the attributes from which it is derived is updated. Triggers can be used to enforce data integrity constraints that cannot be expressed with the PRIMARY KEY, UNIQUE, CHECK, ASSERTION and FOREIGN KEY constructs of SQL because they require manipulating data in more complex ways than can be achieved using keys, checks, assertions, and foreign keys. In Postgres, you will be writing the trigger actions in a procedural programming language that resembles SQL, but adds features such as control structure, variables, etc. This language is called PL/pgSQL. The documentation for PL/pgSQL can be found at <http://www.postgresql.org/docs/8.4/static/plpgsql.html>.

We will also go over it during the help session.

### 4.2 Key and Foreign Key Constraints

#### 4.2.1 Exercises

Listed below are the primary and foreign key requirements for every table in the TPC-H database. Your task is to add these constraints to the database using the appropriate SQL commands. You should do this using `ALTER TABLE` commands only (i.e., not as part of a `CREATE TABLE` command). For example, to add a primary key on an attribute 'thekey' in a table called 'mytable', you would type:

```
ALTER TABLE mytable ADD PRIMARY KEY (thekey)
```

To add a foreign key (calling it `fk1`) from attribute `otherkey` of `mytable`, referencing the primary key attributes of `yourtable`, you would type:

```
ALTER TABLE mytable  
ADD CONSTRAINT fk1 FOREIGN KEY (otherkey) references yourtable
```

For the foreign keys below, do not give any action the default action will be `NO ACTION` which will reject any update that violates the constraint.

Keep your SQL instructions for creating primary and foreign key constraints in a separate file, as you will need to recreate these constraints every time you reset the database for all exercises that follow. The primary and foreign keys to add are as follows:

**Table Customer :**

```
PRIMARY KEY (C_CUSTKEY)  
FOREIGN KEY (C_NATIONKEY) referencing NATION
```

**Table Lineitem :**

```
PRIMARY KEY (L_ORDERKEY,L_LINENUMBER)
FOREIGN KEY (L_PARTKEY,L_SUPPKEY) referencing PARTSUPP
FOREIGN KEY (L_ORDERKEY) referencing ORDERS
```

**Table Nation :**

```
PRIMARY KEY (N_NATIONKEY)
FOREIGN KEY (N_REGIONKEY) referencing REGION
```

**Table Orders :**

```
PRIMARY KEY (O_ORDERKEY)
FOREIGN KEY (O_CUSTKEY) referencing CUSTOMER
```

**Table Part :**

```
PRIMARY KEY (P_PARTKEY)
```

**Table Partsupp :**

```
PRIMARY KEY (PS_PARTKEY,PS_SUPPKEY)
FOREIGN KEY (PS_SUPPKEY) referencing SUPPLIER
FOREIGN KEY (PS_PARTKEY) referencing PART
```

**Table Region :**

```
PRIMARY KEY (R_REGIONKEY)
```

**Table Supplier :**

```
PRIMARY KEY (S_SUPPKEY)
FOREIGN KEY (S_NATIONKEY) referencing NATION
```

For more documentation, see:

- <http://www.postgresql.org/docs/8.4/static/sql-altertable.html> for the ALTER TABLE syntax.
- <http://www.postgresql.org/docs/8.4/static/sql-createtable.html> for table constraint syntax.

## 4.3 Integrity Constraints and Triggers

### 4.3.1 Overview

The data integrity constraints achieved through primary and foreign keys are powerful, but we will see that in complex databases such as the one we are working on, there are more complicated data integrity requirements that require more powerful tools to be expressed. For each data integrity requirement listed below, write the appropriate SQL integrity constraint using ALTER TABLE commands. If (and only if) the constraint cannot be expressed with an SQL integrity constraint, express the constraint with a trigger. We will provide you with details on how to test to see that your integrity constraint is working.

**Note #1** You should work on a fresh database for every problem in this section of the assignment. (I.e., run the reset script from a shell prompt before beginning each new problem). However, you'll need to add the primary and foreign key constraints you formulated in Section 4.2.1 everytime you reset the database, so you are well-advised to keep the SQL instructions that add these constraints in a separate file.

**Note #2** When testing out triggers and their associated functions, Postgres does NOT check the syntax of a function at creation time only at execution time. To test triggers and functions, you will need to drop them each time; drop the trigger first, then the function:



```
DROP TRIGGER foo ON table;
DROP FUNCTION some_function();
```

#### 4.3.2 Exercises

1. If **nation** is updated to change a nationkey for a given nation, nationkeys should be modified across every table in the database.

**To Test Your Solution:** Your solution can be verified by issuing the following queries:

```
SELECT c_custkey, c_nationkey FROM customer WHERE c_nationkey = 99;
SELECT s_suppkey, s_nationkey FROM supplier WHERE s_nationkey = 99;
```

Both queries should return tables with 0 rows.

Next, issue the following update:

```
UPDATE nation SET n_nationkey = 99 WHERE n_name = 'UNITED STATES';
```

Now issue the initial two queries again. If you've implemented the constraint or trigger correctly, these queries should now return 184 and 11 rows respectively. After you've successfully tested your constraint or trigger, reset the database and recreate the primary key and foreign key constraints you formulated in Section 4.2.1.

**Note:** If you need to run this test again, make sure that you reset the database and recreate the primary and foreign key constraints from Section 4.2.1 beforehand. For all exercises in Part 1, you will need to reset the database and recreate the key constraints if you want to test your constraint more than once.

2. If a part's retail price changes, the supply cost for any row in **partsupp** containing this part should change by the same amount. Do not worry about updating prices in **lineitem**.

**To Test Your Solution:** Issue the query below to see the current costs of this part from various suppliers:

```
SELECT ps_supplycost FROM partsupp WHERE ps_partkey = 1;
```

Next, issue the following update query to increase the parts retail price by 10%:

```
UPDATE part SET p_retailprice = p_retailprice * 1.1
WHERE p_partkey = 1;
```

Now, issue the 1st query again to see if the corresponding supply costs in **partsupp** have changed by the appropriate amounts. After you've successfully tested your constraint or trigger, reset the database and recreate the primary key and foreign key constraints you formulated in Section 4.2.1.

3. Customers cannot have more than 14 open orders. Reject any additional orders for customers who already have 14 orders. Be careful, Postgres is case sensitive! An open order is an order with **o\_orderstatus** set to 'O', not 'o'.

**To Test Your Solution:** The customer with **custkey** = 112 already has 13 open orders. Insert another order for this customer:

```
INSERT INTO orders VALUES
(99098, 112, 'O', 99.00, NOW(), '5-LOW', 'Clerk#99', 0, 'IWillPass');
```

Now that this customer has 14 open orders, trying to insert another order should fail. After you've successfully tested your constraint or trigger, reset the database and recreate the primary key and foreign key constraints you formulated in Section 4.2.1.

4. Upon deletion of a customer, delete any orders (and its lineitems) the customer had.

**To Test Your Solution:** Determine the orders/lineitems the customer with **custkey** = 203 has by issuing the query:

```
SELECT l_orderkey, l_linenum
FROM lineitem
WHERE l_orderkey IN (SELECT o_orderkey FROM orders WHERE o_custkey = 203);
```

This should return:

l_orderkey	l_linenum
80996	1
80996	2
65280	1
14945	1
14945	2
14945	3
14945	4
14945	5
14945	6
148071	1
148071	2
148071	3
148071	4
136678	1
136678	2
136678	3
136678	4

(17 rows)

Now issue a deletion query to delete this customer:

```
DELETE FROM customer WHERE c_custkey = 203;
```

The first query should now return 0 rows.

After you've successfully tested your constraint or trigger, **reset** the database and recreate the primary key and foreign key constraints you formulated in Section 4.2.1.

5. A lineitem can have status 'F' (fulfilled), 'O' (open) or 'P' (partial). If a new lineitem is inserted check all lineitems for that order, and update orders. `o_orderstatus` appropriately. Note: If all lineitems have the same (line) status, the order also has this (order) status. Otherwise the order status is 'P'.

**To Test Your Solution:** Issue the following query:

```
SELECT o_orderstatus FROM orders WHERE o_orderkey = 7;
```

This order should have a status of 'O'. Now issue the update:

```
INSERT INTO LINEITEM VALUES
(7, 1, 2, 8, 1, 99.00, 0.00, 0.05, 'N', 'F', NOW(), NOW(), NOW(), 'NONE', 'MAIL', 'No Comment');
```

When you issue the 1st query again, it should report that the order with `o_orderkey = 7` has a status of 'P'. After you've successfully tested your constraint or trigger, reset the database and recreate the primary key and foreign key constraints you formulated in Section 4.2.1.

6. If you delete a region from the database, all nations within that region, and all suppliers located in any of those nations also should be deleted. However, customers from those nations should not be deleted. Any `partsupp` record associated with deleted suppliers should also be deleted, but `lineitem` records referencing those `partsupp` records should not be deleted.

**To Test Your Solution:** Delete the region with regionkey 3 (Europe) from the region table with the following update:

```
DELETE FROM region WHERE r_regionkey = 3;
```

A good way to test your solution for this exercise involves counting how many rows you deleted or affected, although you may wish to verify manually the contents of some of the tables to confirm your solution. Run the following queries:

```
SELECT COUNT(*) FROM nation WHERE n_regionkey = 3;
SELECT COUNT(*) FROM supplier;
SELECT COUNT(*) FROM customer WHERE c_nationkey IS NULL;
SELECT COUNT(*) FROM partsupp;
SELECT COUNT(*) FROM lineitem WHERE l_partkey IS NULL AND l_suppkey IS NULL;
```

If your constraint or trigger was expressed correctly, the queries should return 0, 229, 859, 18407 and 42154 rows respectively. (Before the trigger been executed, the queries would have returned 5, 299, 0, 24021 and 0 rows respectively.)

After you've successfully tested your constraint or trigger, reset the database and recreate the primary key and foreign key constraints you formulated in Section 4.2.1.

7. The table below shows the changes in supply costs for all parts issued by a supplier, when that supplier moves from one region to another. For example, if a supplier moves from Asia to Europe, the costs of parts sold by that supplier should rise by 10%.

From/New	America	Asia	Europe
America	0%	-20%	5%
Asia	20%	0%	10%
Europe	-5%	-10%	0%

Update the costs of parts sold by a supplier according to the rates listed in this table every time a supplier changes its location (**nationkey**).

**To Test Your Solution:** The supplier with **suppkey** = 1 is located in the USA. Issue the following query to determine the total costs of parts sold by this supplier:

```
SELECT SUM(ps_supplycost) FROM partsupp WHERE ps_suppkey = 1;
```

The answer should be \$37704.83. Now, issue the following update to move the supplier's location to France.

```
UPDATE supplier SET s_nationkey = 6 WHERE s_suppkey = 1;
```

Issue the 1st query again. The sum returned should now have changed by the amount indicated by the table above (+5%). Note that Postgres does some rounding in computing the answer to the precision specified by the column type, and for this example the updated supply cost would show \$39590.11. Try this test for other suppliers to test all other percentages in the table.

After you've successfully tested your constraint or trigger, reset the database and recreate the primary key and foreign key constraints you formulated in Section 4.2.1.

## 4.4 To Hand In

For Part 1 of PA2, you will need to submit a file containing the following:

1. The SQL code that adds to the TPC-H database, all of the primary keys and foreign keys described in Section 4.2.1
2. The SQL constraint commands and triggers used to enforce the 7 constraints described in Section 4.3.2.

Each of your primary and foreign key constraints should be written as files under the following directory structure:

```
./sql/keys/[primary / foreign].sql
```

Where all primary key definitions go in the **primary.sql** file and all foreign key definitions go in the **foreign.sql** file. You should be able to pipe them in that order to **PSQL**, which is how they will be graded.

Each of your constraint or trigger used to enforce constraints 1-7 should be written as files under the following directory structure:

```
./sql/triggers/q-[question number].sql.
```

## 5 Java and JDBC

In PA1, you issued queries to Postgres via a command-line-interface (PSQL) and viewed query results as text. In practice, SQL queries are typically embedded in Java code and issued to a DBMS using the [J]ava [D]ata[b]ase [C]onnectivity (JDBC) API, and query results are typically visualized using graphic visualization or form tools. For example, a data analyst might use `gnuplot` to display results as a barchart or scatterplot, a data scientist might use `graphviz` to display data relationships as a directed graph and an online merchant might use `LATEX` to generate a purchase order. For this part of the assignment, you will be writing Java code with embedded SQL queries that will be issued to Postgres using JDBC, and then processing query results for subsequent input to `gnuplot` (5.1), `graphviz` (5.2) and `LATEX` (5.3).

### 5.1 Task #1 - Visualizing Supplier Sales by Region

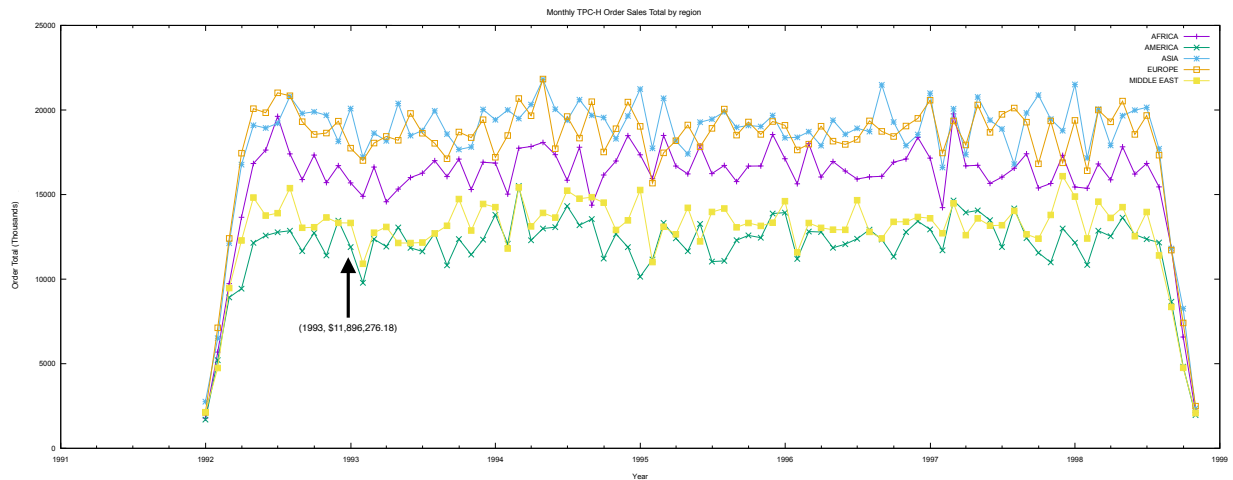


Figure 2: Task #1 output using `gnuplot`

For this task, you will write Java code that issues SQL queries to the TPC-H database to calculate supplier sales totals by region over time and then displays the results as a *scatter plot* using `gnuplot`. Figure 2 shows the result your code should generate. Observe that the x-axis of this graph shows each year from 1992-1999, and the y-axis of this graph shows the total income from sales of suppliers in each of 6 continents (differentiated by plot line color) for each month of these 8 years. For example, the indicated data point in this Figure shows that the sum of all sales from suppliers in America from selling items shipped in January of 1993 was \$11,896,276.18.

**HINT:** The granularity of this result is monthly. The `EXTRACT` function retrieves subfields such as year or hour from date/time values:<sup>4</sup>

`EXTRACT(field FROM source)`

Where *source* must be a value expression (such as an attribute) of type timestamp, time, or interval, and *field* is an identifier or string that selects what field to extract from the source value. The return type is of type double precision. For example, in a table *person* with a column *dob* of type `date`:

```
SELECT EXTRACT(DAY from dob) FROM person; → [day of birth]
SELECT EXTRACT(MONTH from dob) FROM person; → [month of birth]
SELECT EXTRACT(YEAR from dob) FROM person; → [year of birth]
```

<sup>4</sup>Source: <http://www.postgresql.org/docs/9.4/static/functions-datetime.html>

## 5.2 Task #2 - Visualizing Trade Relationships by Region




Figure 3: Task #2 output using `graphviz`

For this task, you will write Java code that issues SQL queries to the TPC-H database to calculate the total sales between suppliers and customers grouped by the regions where the suppliers and customers are located. The results should then be displayed as an annotated *directed graph* using `graphviz`. Figure 3 shows the graph that your code should generate. Note that each node in the graph represents a region (continent), and the directed edge from a region  $r_1$  to another region  $r_2$  ( $r_1 \rightarrow r_2$ ) is annotated with the total sum of purchases made by customers located in region  $r_1$  from suppliers located in region  $r_2$ . For example, the total purchases by customers in “Asia” from suppliers in “America” is \$190M. Note that total sales should be computed in both directions for each pair of regions (e.g., you need to determine **both** the total purchases by Asian customers from American suppliers and the total purchases from American customers from Asian suppliers) and should also include total sales within a given region (e.g., the total purchases by American customers from American suppliers).

### 5.3 Task #3 - Generating Purchase Orders

Brandeis University

415 South St.

(781) 736-2000

Waltham, MA 02453

Date:

10/24/2015

Order #:

4

Part #	Quantity	Unit Price	Amount
Supplier: 106			
1223	95	\$253.55	\$24087.25
Subtotal			\$24087.25
Supplier: 108			
4801	37	\$344.32	\$12739.84
Subtotal			\$12739.84
Supplier: 109			
1822	50	\$53.60	\$2680.00
3098	12	\$116.39	\$1396.68
3229	59	\$392.26	\$23143.34
4714	4	\$14.38	\$57.52
Subtotal			\$27277.54
Supplier: 11			
1206	15	\$82.20	\$1233.00
679	4	\$177.54	\$710.16
Subtotal			\$1943.16
Total			\$66047.79
# of Suppliers			4

Figure 4: Task #3 output using L<sup>A</sup>T<sub>E</sub>X

For this task, you will write Java code that issues SQL queries to the TPC-H database to determine minimum prices for each of a specified set of items, and sends the results via the provided Java API to L<sup>A</sup>T<sub>E</sub>X to generate a purchase order that orders each item from its cheapest supplier.

We have provided you with a test file (`./txt/orders.txt`) that specifies sets of items for 10 separate orders. Your code should process this file to generate 10 separate purchase orders that show each item in the given order and its price from its cheapest supplier. Figure 4 shows an example purchase order that you should generate for the 4<sup>th</sup> order in the provided file. Note that your generated purchase order should be subdivided into suborders representing the portions of the order being submitted to each supplier. This is illustrated in Figure 4 of purchase order #4, containing 4 suborders: one to supplier 106 for part 1223, one to supplier 108 for part 4801, one to supplier 109 for parts 1822, 3098, 3229, and 4714, and one to supplier 11 for parts 1206 and 679.

File `./txt/orders.txt` consists of a single line for each order with format,

(order number):(part quantity)x(part number), ...

For example, order number #957 for 10 units of part 195 and 30 units of part 227 is shown as:

957:10x195,30x227

The output of your program should be L<sup>A</sup>T<sub>E</sub>X code producing a document consisting of one purchase order (such as that shown in Figure 4) for each order specified in `./txt/orders.txt`.

**IMPORTANT:** Your code for this task should work for any number of orders and any number of items per order. Be forewarned that we will test your code with a different version of `./txt/orders.txt` than the one being provided to you to complete this assignment.

## 5.4 To Submit

Once you are done with the 3 tasks above, please submit the following:

- All your source code with comments in `./src`
- Your `gnuplot`, `graphviz`, and `LATEX` source for your output located in `./plot`, `./gv`, and `./tex` respectively
- Your plot, graph, and purchase order located in `./viz`

## 6 Submission

Submitting both parts 1 and 2 of this assignment is easier than ever. First, set the proper environment variables:

```
$ export FIRSTNAME=[your first name as it appears on LATTE]
$ export LASTNAME=[your last name as it appears on LATTE]
```

Then simply type the following command to package all relevant files into a submittable file:

```
$ make submit
```

A file with the name `[your last name]_[your first name]_127b_pa2.tar.gz` will be created in the current directory.

Then you can use secure copy to transfer the file from the server to your machine. Simply execute the following command from your local machine:

```
$ scp [your username]@[cs public machine name].cs.brandies.edu: /[pa2]/*_*_127b_pa2.tar.gz .
```

The file should then be transferred to the current directory.

Upload the file to LATTE.